

Xingfan Xia

8502 134th Ct NE, Redmond, WA 98052
xingfanxia@gmail.com • +1 (507) 403-1689 • <https://ax0x.ai>

PROFILE

ML/AI engineer and startup CTO with 7+ years building production ML systems — from on-device models at Apple to multi-agent AI orchestration at scale. Currently shipping AI-powered automation via **Agentic AI Development** (Claude Code, 80B+ tokens). Building evaluation pipelines and multi-agent orchestration methodology for enterprise AI.

WORK EXPERIENCE

Compute Labs, Remote, Redmond, WA

- Co-Founder and CTO Aug 2024 – Present
 - **Architecting multi-agent AI systems for GPU infrastructure underwriting, building LLM-powered automation that reduced deal evaluation from 2 weeks to under 10 minutes.**
 - Architected multi-model AI orchestration layer (Claude, Gemini, OpenAI) powering 11 parallel domain-expert sub-agents for automated deal evaluation, with three-layer pipeline: Gemini extraction, Claude Opus analysis, and multi-model ensemble for report generation.
 - Built state machine workflow engine orchestrating AI agent pipelines processing \$1.3B+ in GPU deals, with prompt experimentation system for continuous model optimization.
 - Built GPU market intelligence platform using multi-model LLM extraction (GPT/Claude/Gemini) to process 3,200+ data points from 45+ providers, with 4-stage ML pipeline including adaptive outlier detection and financial depreciation modeling.

Amazon Web Service – Athena, Redmond, WA

- Senior Engineer - ML Data Platform Apr 2022 – Jun 2024
 - **Building petabyte-scale query engine powering ML training data pipelines and feature stores, enabling data scientists to prepare, transform, and materialize datasets for model training at AWS scale.**
 - Designed and shipped feature materialization system enabling ML teams to version, transform, and snapshot training datasets at petabyte scale with automated schema evolution for iterative model development. Released at AWS RE:INVENT 2022.
 - Built high-throughput ML data processing pipeline optimizing feature extraction workflows, improving end-to-end throughput by 30% and reducing data pipeline failures by 60%.

Apple, Cupertino, CA

- Machine Learning Engineer - Siri AI Platform Aug 2020 – Apr 2022
 - **Building on-device ML models and AI-powered features for Siri proactive assistant across iOS, macOS, watchOS, and tvOS.**
 - Implemented on-device user profiling framework with multiple ML models predicting user interests and preferences, serving as upstream signal provider for Apple Maps and Apple News recommendation systems.
 - Overhauled topic prediction model using transformer-based BERT architecture, increasing prediction accuracy by 14% and user engagement by 11%. Optimized inference pipeline to reduce on-device latency by 12%.

Airbnb, San Francisco, CA

- Machine Learning Engineer - Trust & Safety Jan 2019 – Aug 2020
 - **Building ML-powered fraud detection and trust & safety systems using XGBoost, logistic regression, and real-time inference engines.**
 - Developed ML models (XGBoost and logistic regression) for fake review detection, cutting fraudulent reviews to 1% of peak volume and preventing \$2M in losses. Deployed with optimized inference mechanism improving performance by 30%.
 - Architected real-time ML inference microservice with Kafka streams and Redis caching for sub-100ms risk signals across multiple downstream consumers.

EDUCATION

Carleton College, Northfield, Minnesota

- B.A. in Computer Science Graduated Jun 2018

TECHNICAL SKILLS

- Programming Languages: Proficiency in Python, Java, Golang, TypeScript, Scala, Objective-C
- AI & Machine Learning: LLM APIs (Claude, GPT, Gemini), Multi-Agent Systems, Prompt Engineering, PyTorch, Scikit-learn, XGBoost, BERT, On-Device ML, MLOps, Claude Code, Cursor
- Infrastructure & Data: FastAPI, Next.js, React, GCP Cloud Run, Spark, Iceberg, Kafka, Airflow, AWS (Athena, S3, Lambda, EC2)